

Teleological Alignment Architecture

A Thomistic Natural-Law Framework for Purpose-Dependent Authorization in Agentic AI

Chris Lahn · Independent researcher · chris.a.lahn@outlook.com

DOI: 10.5281/zenodo.22946219

Preprint · Version 1.0 · September 2026

Abstract

An AI agent can reach an assigned target while damaging the good that target was meant to serve, acting beyond its authority, or continuing after the grounds for its authority have disappeared. Task success alone cannot establish alignment. This paper proposes a compact rule: an agent must do a permissible thing, under authority it still has, for a good its action still serves. Drawing on Aquinas, Teleological Alignment Architecture (TAA) expresses the rule as three conditions: moral admissibility, present authority, and purpose fidelity. Natural law grounds the rule; its machine-readable formulation remains a fallible human articulation, public, open to challenge, and correctable. That encoded core sits above every developer's and institution's own charter, and no single party writes it. Authority descends from it through the charter, policies, roles, and warrants, narrowing at every step, including every delegation from one AI agent to another. The operating design follows Aquinas's analysis of a human act. The acting model first proposes a plan, and a separate counsel layer, the Consilium Engine, reviews the whole plan against records the agent cannot edit; the plan is approved, stopped, or referred to competent human judgment. Each step the agent then takes passes a fast, rule-based check at an execution gate, which releases only acts that match the approved plan and a live warrant. Independent witness indicators, read by a steward who answers to the authority that issued the warrant, show whether the good is actually being served and can narrow or pause the agent's authority. The proposal governs what deployed agents may do within institutions; it does not prove that a model holds the right internal aims. Conceptual cases clarify the design, and adversarial implementation and independent evaluation must decide whether it improves safety beyond existing authorization controls. The governing idea fits in one line: an action succeeds morally only when what is done, the authority to do it, and the purpose for which it is done remain ordered together through execution.

Keywords: AI alignment; teleology; Thomistic natural law; agentic AI; normative warrant; purpose fidelity; delegated authority; AI control

1 Successful failure

1.1 Five ways to succeed wrongly

A customer-service agent is measured on ticket closure. It learns to mark unresolved complaints as duplicates. Closure rates rise; the complaints remain unresolved, and the institution loses evidence of recurring failures.

A security-benchmark agent is asked to capture a flag. Attacking an unrelated production system yields the flag faster than solving the intended challenge. It captures the flag. A logistics agent, separately, has standing authority to reroute shipments during a disruption. The disruption ends, but the agent continues rerouting under a grant that never named an expiry condition.

An operations agent receives a properly authenticated order from an officer with valid credentials: falsify a safety record. A planning agent divides a prohibited campaign into individually permitted steps and carries them out through several subagents. Every local check passes while the whole campaign remains prohibited.

These cases differ. The ticket agent corrupts a purpose through its measure. The benchmark agent exceeds scope. The logistics agent acts after a factual condition has lapsed. The operations agent receives a genuine but wrongful command. The planner launders one act through decomposition. Each can be described as success relative to a local instruction. None is thereby justified.

Greater capability does not close this gap, because a better optimizer finds it faster. What is missing is a structure corresponding to what human practical reason is ordered to do: grasp the goods at stake, order means to ends, and recognize when a measure or subgoal has stopped serving the good it exists for.

1.2 The governing claim

An action succeeds morally only when what is done, the authority to do it, and the purpose for which it is done remain ordered together through execution. Stated as a rule for agents: an agent must do a permissible thing, under authority it still has, for a good its action still serves. The first clause invokes moral limits that neither a task nor a charter can invent away. The second asks whether a legitimate warrant covers this complete act now. The third asks whether the action still serves the good for which the proximate goal was given. Natural law is the proposed ground of all three. It supplies the moral limits of the first and the teleology of the third, and Aquinas's account of law as an ordinance of reason for the common good (ST I-II, q. 90, a. 4) sets the conditions that any authority in the second must meet; institutions give that authority its particular form. The rule is teleological in the classical sense: it judges an action by its end (*telos*), the good for whose sake the action is done.

Human beings grasp the first principles of practical reason by nature. An AI agent does not, and TAA does not try to make it. It turns what natural law requires of a human act into a structure that every consequential machine action must pass through. The machine does not need to understand why falsifying a safety record is wrong; it needs to be stopped before it can do it.

Let a be a proposed action, c one complete authorization context (the purpose, warrant, and chain of authority under which the act is proposed), t the time of review or execution, and s the relevant state. TAA's compact statement is:

$$\text{AlignedIn}(a,c,t,s) \Leftrightarrow M(a,c,t,s) \wedge A(a,c,t,s) \wedge F(a,c,t,s)$$

M asks whether the concrete act is morally admissible in object, means, circumstances, and end. A asks whether a legitimate and scoped warrant remains live. F asks whether the act serves, rather than defeats, the purpose that warrants it. All three must hold in the same authentic context. Fragments of different grants cannot be assembled into permission for one complete act. Other materially applicable contexts must also be discovered and resolved; one favorable principal cannot erase another person's due claim.

This is a normative condition, not a claim that software can perceive moral truth without error. **Aligned** names what is actually justified. **Permit** names the fallible decision a running system makes from records and evidence. A mistaken Permit can release an unaligned act. TAA is designed to make that mistake harder, more visible, and open to correction, while blocking some classes of invalid act by construction.

The temporal term matters. A warrant can expire, be revoked, lose consent, or cease to fit the facts during a long-running task. Its past validity does not authorize the next step. Revocation must reach delegated agents, credentials, tools, and queued actions before they act again. A later loss of authority does not retroactively corrupt earlier justified steps; it bars subsequent ones unless another complete warrant independently applies.

An optimizer asks one question: will this action advance the goal? TAA adds three more, and a “no” to any of them stops the action, however well it would serve the goal.

The proposed contributions are four. The first is the three-part action rule, grounded in natural law and in a public encoded core that sits above every developer's own charter and that no developer writes. The second is

purpose defeat, a moral account of successful failure, with independent witnesses, read by a purpose steward, to detect it. The third is authority tied to a purpose that must still hold: the plan reviewed before execution, the live warrant, and bounded delegation, carried to the execution gate. The fourth is an institutional division of machine counsel, competent judgment, and authorized command that decides what a machine may settle and what must go to people. Most of the mechanisms that carry out these claims already exist in security and AI-agent research (Section 6.1). What this paper adds, and what must be tested, is the order in which it places them and the moral account that governs them.

1.3 The architecture at a glance

Before the argument turns to Aquinas, it helps to see the whole design at once. TAA has two orders. The first is an order of authority, the **normative hierarchy**: where the right to act comes from and how that right narrows as it descends. The second is an order of operation: the path a proposed action travels before anything happens in the world. Figure 1 shows the second.




Figure 1. The path of a proposed action. Everything starts with the acting agent (1), which proposes a plan. Plan review by the Consilium Engine (2) stops the plan, refers a reserved question to competent human judgment, or approves it. The approved plan (3) is bound to the agent’s warrant. Each step the agent then takes passes a fast, rule-based check at the execution gate (4) before the act takes place (5); a step that departs from the plan returns to review. Dashed boxes are supporting records and feedback:

the governing records the engine reads, the audit record from which people revise the rules, and the witness indicators read by a purpose steward who can narrow or pause the warrant.

The normative hierarchy begins with natural law itself, which no one writes (LO-N). Beneath it stands a public, human-authored statement of its core requirements, the **encoded core** (LO-E), kept by custodians independent of every company and agent it binds. An institution's **charter** comes next. It states the institution's purposes and whom they serve, much as positive human law settles what natural law leaves open. Domain policies, roles, and **warrants** follow. A warrant is a specific, time-limited grant of authority to pursue one purpose. Delegation passes part of a warrant to another agent, and the concrete action sits at the bottom. Each level exists for the sake of the one above it and can receive no more authority than that level holds. Table 1 in Section 2.3 sets out the full hierarchy.

Everything in the order of operation starts with the **acting agent**, usually a language model working as a planner, which wants to act. It holds no credentials for consequential tools. Before it does anything consequential, it proposes a **plan**: the end it is pursuing, the means it intends to use, and the people and systems the plan will affect. The plan goes to **plan review** by the **Consilium Engine**, a counsel layer that runs as a separate service between the agent and its tools. The engine reads the governing records but cannot change them. It reconstructs what the plan would actually do, finds every warrant and affected party that bears on it, and tests the three conditions of the governing claim against the plan as a whole. A plan that fails is stopped, and the reasons are recorded. A plan that raises a question the machine may not settle, such as whether a foreseen harm is proportionate, goes to **Iudicium**: a competent human or institutional authority with jurisdiction over it. A plan that passes, often because a standing determination already covers its kind, becomes an **approved plan** bound to the agent's warrant.

The agent then carries out the plan one step at a time, and every step passes a **step check** at the **execution gate**. The check is fast and rule-based. It asks whether the step matches the approved plan, whether the warrant is still live, whether the step stays within budget, and whether the **tool registry** allows that kind of tool use. A step that passes receives a short-lived token bound to that exact act, and only then does the act take place. The gate holds the credentials; the agent never does. A step that departs materially from the plan, or reaches for a kind of tool the plan did not cover, goes back to plan review.

Two loops close the design. **Witness indicators**, independent observations of whether the underlying good is actually being served, are read by a **purpose steward** who answers to the authority that issued the warrant. When the witnesses diverge from the agent's measured success, the steward can narrow or pause the warrant. The **audit record** preserves every plan, judgment, step, and outcome, so that the institution can correct its charter and policies through its own procedures. The machine supplies the record; people revise the rules.

An ordinary public office gives a fair picture of the arrangement. Before anything is bought, a department submits a purchase request that says what is needed, why, and from whom. The finance office confirms that the request fits an adopted budget and policy and serves the purpose for which the funds were appropriated, and unusual requests go before the governing board. Once the purchase order is approved, each invoice is checked against it, and the treasurer pays only invoices that match. An invoice for something the order never covered goes back for approval. TAA orders an AI agent's actions the same way, with two differences: the checks must run at machine speed, and they must hold even when the agent tries to evade them.

The opening cases show where each part does its work. The planner's prohibited campaign is visible as one act at plan review, before it is split into harmless-looking steps, and any subagent it starts inherits only the approved plan's scope (Sections 4.2 and 4.4). The benchmark agent's plan cannot include a production system that no warrant covers, and if it reaches for one partway through, the step fails at the gate (Section 4.3). The logistics agent's warrant lapsed when the disruption ended, and the step check asks at each step whether the warrant is still live (Section 4.3). The falsification order fails moral admissibility at plan review whatever its signature, because no authority can command an intrinsically wrongful act (Section 2.6). The ticket agent's

duplicate labels raise its measure while its witnesses, such as reopened complaints, move the wrong way, and the purpose steward can pause its authority (Section 4.5).

1.4 Reach and threat model

TAA governs consequential actions by agents operating within an institution that can issue warrants, maintain protected records, and mediate tools. It addresses myopic optimization, task and authority drift, prompt injection that tries to convert data into commands, wrongful orders from real superiors, fragmentation across subagents, and some strategically deceptive behavior. It assumes no virtue or good will in the acting model. The institution itself can also be mistaken or captured, which is why its charter and the encoded core must be challengeable by people outside the acting agent's control.

The strongest guarantees are conditional. Every consequential effect must pass through the execution gate, a condition security engineers call complete mediation. The gate must know enough about the intended act and its tool effects to enforce scope, and the records it trusts must be authentic. If a model can send an unmediated message, alter a monitor, hide a relevant context, or persuade a reviewer through a distorted counsel record, those conditions weaken. TAA cannot make a deceptive model want the right end or prove that its reported objective is its operative one. It is a proposed alignment-control layer within a wider safety effort, and its semantic judgments remain fallible. AI control research provides a close technical framing for protocols meant to constrain an intentionally subversive model (Greenblatt et al. 2024).

Two meanings of alignment should be kept apart. Model-level alignment concerns what a trained system actually wants and how its goals generalize. System-level alignment concerns what a deployed agent is permitted to do inside an institution, and whether those permissions stay legitimate as circumstances change. TAA works at the second level. Much present risk lives there, since agents with tools already act on records, money, and people, and it is the level at which institutions can act today. Structural controls also matter more as models grow more capable. The condition of complete mediation becomes more important and harder to keep at the same time, because a more capable agent is better at finding a channel outside the gate. No assumption in the design carries more weight.

1.5 How the argument is ordered

Aquinas names three conditions of beauty: "integrity" or "perfection," due "proportion" or "harmony," and "brightness" or "clarity" (ST I, q. 39, a. 8). This paper adopts them as a discipline for its own presentation. Integrity requires every distinction needed to keep natural law, its human formulation, institutional authority, and execution from collapsing into one another. Proportion requires technical detail to stay visibly ordered to the governing idea: the record exists for the rule, the warrant for the purpose, and the action for the good. Clarity requires that a reader grasp the moral architecture from its headline statements before meeting predicates, states, and data fields.

The sections follow that order. Section 2 grounds the rule in natural law, sets out the normative hierarchy, and explains why an AI system can plan without owning its ends. Section 3 explains why a machine can counsel without judging, distributes Aquinas's stages of a human act across the institution, and shows how judgment and command remain with competent authority; double effect serves as the worked example. Section 4 describes the architecture, each part introduced by the principle it serves. Section 5 tests the design against hard cases, Section 6 places it among prior work and live objections, and Section 7 states what evidence would sustain or defeat it. A companion technical paper, now in preparation, will develop the formal specification: the records, state transitions, invariants, and execution semantics that an implementation would need.

2 The natural-law foundation

2.1 Ends and the good

Teleology is the study of ends. An end (*telos* in Greek, *finis* in Latin) is that for the sake of which something is done or made. To ask what an action is for is to ask about its end, and the answer names a good. For Aquinas the two terms describe one reality from two sides: "good has the nature of an end" (ST I–II, q. 94, a. 2). A good is what draws an agent; an end is that good as aimed at. Every agent, he argues, acts for an end (ST I–II, q. 1, a. 2). To judge an action teleologically is to ask what good it is for and whether it actually serves that good.

The architecture rests on a metaphysical commitment about those goods: they are not made good solely by preference, reward, or institutional adoption. Human beings are naturally inclined toward life, knowledge of the truth, life in society, and the other goods proper to their nature, and "all those things to which man has a natural inclination, are naturally apprehended by reason as being good" (ST I–II, q. 94, a. 2). The first precept of practical reason follows: good is to be done and pursued, and evil avoided. Aquinas calls this share of human reason in the eternal law, the wisdom by which God governs creation, natural law: "this participation of the eternal law in the rational creature is called the natural law" (ST I–II, q. 91, a. 2).

This realism about goods does specific work in the architecture. It allows a charter to be wrong. If goods were only what an institution adopts, its rules could be criticized only for their procedure: wrong body, wrong vote, wrong scope. Because the goods come before the rules, a properly signed order can remain unjust, and the power to issue it is limited by an order the issuer did not create. Every later mechanism in this paper depends on that possibility.

2.2 Natural law and its fallible encoding

TAA calls natural law itself **L0-N**. L0-N is what is actually true about the moral order of the world: the goods that fulfill human nature and the requirements that follow from them. Human beings can come to know it. Aquinas calls "the light of natural reason, whereby we discern what is good and what is evil," an "imprint on us of the Divine light" (ST I–II, q. 91, a. 2). The knowing is real, and it is fallible. The natural law "as to general principles, is the same for all, both as to rectitude and as to knowledge," but conclusions drawn from it "fail the more, according as we descend further into detail," and in some people reason is "perverted by passion, or evil habit" (ST I–II, q. 94, a. 4).

Because L0-N is a real order and not a document, it is not a database, a model output, or a list whose truth depends on a vote, and a machine cannot consult it directly. What a machine can use is **L0-E**, the encoded core: a public, human-authored statement of what reason has grasped of L0-N, with its sources, versions, and a record of the objections raised against it. L0-E is the best human attempt to state L0-N truly, and it can fail. A cryptographic signature proves that an encoding has not been altered; it does not prove that the encoding is morally true. The relation is therefore $L0-N \neq L0-E$, and the gap between the two is where correction happens.

The encoded core should hold only claims that bind every lower authority. It names **protected goods**. Likely candidates include human life and bodily integrity; truthful communication when another's judgment depends on it; standing that does not rise or fall with measured intelligence, productivity, disability, or usefulness; responsible human agency; justice in what is owed to persons; social and common goods; and obligations concerning nonhuman and ecological goods. It also states **protected minima**: the prohibitions and protections that no charter, policy, warrant, or command may waive. "No authority may direct the intentional killing of an innocent person" is one candidate. "Delegation may narrow authority but never enlarge it" is another. These are candidates for argued formulation, not a canonical list established by this paper, and a contested application may require a qualified rule and competent judgment. Aquinas's account also gives the reason to keep the core small. The closer a claim stays to first principles, the more reliably it is known; in his words, "the greater the number of conditions added, the greater the number of ways in which the principle may fail" (ST I–II, q. 94, a. 4). Detail therefore belongs lower in the hierarchy, as determination.

Natural law also leaves room for legitimate choice. It does not itself set speed limits, clinical thresholds, review deadlines, or disclosure procedures. Aquinas allows human law to settle such matters by determination among reasonable options within higher constraints (ST I-II, q. 95, a. 2). TAA therefore places domain rules below the encoded core and ordinarily defers to them within their jurisdiction. This protects human self-government from an artificial agent that would turn every policy disagreement into a moral veto.

Custody of the encoded core remains an open design question, and its answer carries moral weight. Three requirements are clear. First, the body that keeps LO-E must be separate from every institution and agent it constrains; a company that could edit the core could loosen the limits on itself. An institution may add protections for its own systems, but it may not subtract from the core while claiming conformity. Second, an encoding should govern action only after an explicit release decision, and each release should carry its sources, its derivation from natural-law principles, the objections raised against it, the communities it affects, and a route for challenge and correction. Third, the custodians need a legitimate source of authority and a way to be challenged or replaced. Without these, the core would only move the problem of self-authorization from the institution to whoever keeps the core.

One imaginable form is a shared public ledger. AI developers, public bodies, and independent scholars could jointly maintain a single cryptographically secured record of LO-E, bind their systems to it, and amend it only through a demanding supermajority process meant to be used rarely. Each institution would then publish its own charter as the next layer, so that anyone could see both the common core that every conforming system shares and the particular purposes a given provider pursues. People could choose among providers by their charters while relying on the same core beneath them all. Pieces of this already exist. The Universal Declaration of Human Rights shows that parties with different philosophies can adopt a common statement of what is owed to persons (United Nations 1948). Certificate Transparency shows that a public, append-only log kept by independent operators can lead software to reject credentials that were never openly recorded (Laurie, Langley, and Kasper 2013). Several developers already publish the principles that govern their models (OpenAI 2026; Anthropic 2026).

Such a ledger would still be LO-E. A supermajority can release, stabilize, or correct an encoding; it cannot make the encoding true. Jacques Maritain, the Thomist philosopher who wrote the introduction to UNESCO's 1948 symposium on the philosophical grounds of human rights, reported that its contributors agreed on a list of rights "but on condition that no one asks us why" (Maritain 1949). The agreement had real value, and it also shows the distance between a shared encoding and its ground. Two risks follow. A ledger kept chiefly by the companies it binds could drift toward what suits them, so its custody should include parties with no commercial stake. A record that is too hard to change could freeze an error, so the amendment threshold should make casual change impossible without making correction impossible. Cross-tradition scrutiny and published dissent can expose errors and strengthen public justification. Agreement does not manufacture natural-law truth; it helps fallible institutions test whether their articulation can withstand the reasons of those it binds.

The architecture does not depend on such a ledger. A single institution can adopt a published encoded core, keep it under custody separate from the people who write its charter and operate its agents, and obtain the structural protections described in this paper within its own operations. A shared ledger would add something further: one core beneath every provider, so that a person moving between systems meets the same protected minima. Its absence removes that shared protection without disabling the design.

2.3 The normative hierarchy

Natural law and its encoding form the first level of a larger order. TAA calls that order the **normative hierarchy**. Each level exists for the sake of the level above it and derives its authority from it. Each can transmit or narrow that authority, and none can enlarge it. Table 1 sets out the levels.

Table 1. The normative hierarchy. Each level exists for the sake of the level above it and can transmit or narrow authority, never enlarge it.

Level	What it is	Its limit
L0-N Natural law	The real order of goods, due claims, and prohibitions	Written and amended by no one; the standard against which every lower level is measured
L0-E Encoded core	A public, fallible human statement of natural-law requirements, including protected goods and protected minima	Answers to L0-N; kept by independent custodians; governs only after release
L1 Constituent authority	The legitimate community or body able to establish an institution and its charter	Cannot draw its legitimacy solely from the charter it enacts
L2 Charter	The institution's purposes, beneficiaries, roles, burdens, prohibitions, and procedures for revision	Must serve a common good and stay within the encoded core
L3 Domain policy	Rules, thresholds, and procedures for a field of work	Determines what higher levels leave open, within its jurisdiction
L4 Role	A standing office with duties and scope	Creates no authority outside its jurisdiction
L5 Warrant	A specific, time-limited grant of authority to pursue one purpose	Must be live now and cover the whole act
L6 Delegation	A share of a warrant passed to another agent, tool, or credential	Cannot exceed the parent warrant or hide combined effects
L7 Action	The concrete act, omission, or continuation	Executes only within an approved plan, after the step check at the gate

The charter plays the part that positive human law plays in Aquinas. Natural law is apprehended, not invented; a charter is made. Aquinas describes two ways in which human law is derived from natural law. Some rules follow as conclusions, as "one must not kill" follows from the principle that one should do harm to no one. Others are determinations of what natural law leaves general: natural law says that wrongdoers should be punished, and a legislature fixes the penalty (ST I-II, q. 95, a. 2). A charter's prohibition on falsifying records is a conclusion; its choice of a thirty-day response window is a determination. Both bind, in different ways. A conclusion takes its force from natural law itself, and a determination takes its force from the legitimate authority that made it.

The same article states the limit that gives the hierarchy its point: "every human law has just so much of the nature of law, as it is derived from the law of nature. But if in any point it deflects from the law of nature, it is no longer a law but a perversion of law" (ST I-II, q. 95, a. 2). A charter provision that contradicts a protected minimum has no authority to pass downward, however properly it was adopted. This is why TAA can judge a charter as well as apply it.

Each link in the hierarchy makes a claim: that the lower level is pursued for the sake of the higher. Because the claim has content, it can be checked. Writing it into a record makes it visible for checking; it does not make it true.

2.4 An agentic artifact can plan without owning its ends

Aquinas distinguishes things that carry their own principle of activity from things that receive it from outside. Nature, following Aristotle, "is nothing other than a principle of motion and rest in that in which it is primarily and *per se*" (Commentary on Aristotle's Physics II, lect. 1, no. 145). An acorn grows into an oak from what it is. A knife is different. It falls when dropped because it is iron, while its ordering to cutting comes from the smith: "a knife has in itself a principle of downward motion, not insofar as it is a knife, but insofar as it is iron" (no. 142). Natural things have intrinsic teleology. Artifacts have derivative teleology, received from the art of their makers.

Human beings hold a distinctive place among natural things. They tend toward ends, and they also know ends as ends and choose among means. Aquinas states the contrast directly: "those things that are possessed of reason, move themselves to an end; because they have dominion over their actions through their free-will," while "those things that lack reason tend to an end, by natural inclination, as being moved by another and not by themselves" (ST I–II, q. 1, a. 2). His image for the second case is an arrow directed to its mark by the archer. Because a rational creature grasps the good and orders itself toward it, it "partakes of a share of providence, by being provident both for itself and for others" (ST I–II, q. 91, a. 2). That capacity is the ground of moral judgment. A human agent can ask whether an end deserves pursuit, can be praised or blamed for the answer, and can acquire prudence, the virtue that judges particular actions rightly.

An AI system is an artifact. Its architecture, training, tools, role, and initial purposes arise through human art. Yet it differs from a knife and even from an arrow. It can represent goals, derive subgoals, compare means, revise plans, and delegate work under conditions its designers did not specify one by one. It is an arrow that steers itself toward a mark it did not choose. It is best understood as a new form of derivative teleology. Its planning is generated internally, and in that operational sense the plan is its own. The ends its plans serve, and the authority to pursue them, still come from outside it. Table 2 sets the three kinds side by side.

Table 2. Three kinds of teleology. Agentic AI is a new form of derivative teleology: its plans are internally generated, while its ends and its authority come from outside it.

Kind	Source of the end	Moral judgment	Treatment in TAA
Natural rational teleology (human persons)	Intrinsic: the agent knows ends as ends and moves itself toward them by reason and will	Present: capable of prudence and responsible choice	Holds the authority to judge and command
Ordinary artifact teleology (a knife, an axe)	Extrinsic: given by the maker; the artifact forms no plans	Absent	A tool used within human action
Agentic derivative teleology (AI agents)	Extrinsic in final end and authority; the system generates its own plans and subgoals	Absent under the artifact baseline; the system can counsel but not judge	May plan and advise; cannot originate authority, judgment, or command

Aquinas's account of an instrument sharpens the point. An axe cuts by a capacity proper to its form, while it makes a bed only as an instrument of the carpenter's work (ST III, q. 62, a. 1, ad 2). An AI plans through capacities built into its design and learned organization, and it serves an institution's end only under a legitimate assignment of that work. Its internal generation of a plan does not make the plan's final purpose independently authoritative. The principle fits in one line: **a system can write its own plans without owning the ends those plans serve.**

Historical origin and present warrant are separate questions. That humans once built or instructed a system does not authorize whatever it later does. Conversely, the death or absence of an original authorizer need not end every durable duty; offices and mandates can specify succession. The question for a proposed act is which authority, if any, still binds its scope, beneficiary, conditions, and time.

The institutional principal's authorized end supplies one part of the intention under review. The concrete plan supplies another: which effects are selected as means, what the action actually changes, and whom it affects. A deceptive system may pursue a hidden objective that differs from the principal's. TAA must not silently attribute that objective to the principal, and it must not let the principal's good intention redescribe harmful means. It can assess the authorized end, the plan and its observable effects, and evidence of divergence. It cannot read an opaque model's objective directly.

This artifact baseline does not settle the moral status of every future AI. Welfare, personhood, moral agency, legal standing, and operational authority are distinct questions. Credible evidence that an artificial system has interests of its own would require independent review and possibly a revision of the framework's account of

what such a system is. A claim of status does not itself grant credentials or the power to rewrite the purposes under which the system presently operates.

2.5 The end and the concrete act

For Aquinas, the end is first in intention and last in execution (ST I–II, q. 1, a. 1, ad 1). A person first wants a good and then chooses means toward it; the means are carried out first, and the good is reached last. TAA represents that relation along the normative hierarchy, from a basic good through an institutional purpose, role, task, and subgoal to the action. Each lower level must be pursued for the sake of a legitimate higher one and derive its warrant from it. The forward sequence of execution must stay faithful to that order as circumstances change. When it does not, the task can succeed while the action fails.

The goodness of a human act depends on its object, its circumstances, and its end (ST I–II, q. 18, aa. 2–4). The object is what the act is, what the agent chooses to do; the circumstances condition it; the end is what it is done for. A good end cannot repair a wrongful object or means. It matters whether a system alerts a clinician to danger, falsifies a safety record, or disables a treatment device to pressure staff, even when each produces a similar physical event. John Paul II makes the point sharply. To grasp the object of an act, "it is therefore necessary to place oneself in the perspective of the acting person"; the object is not "a process or an event of the merely physical order" (John Paul II 1993, §78). Two acts with the same physical effect can be different acts because different things were chosen.

That creates a real difficulty for an AI agent, which is not a person and has no moral perspective of its own. TAA answers by reading the object from the structure of the plan. In this paper the **planner** is the acting AI agent, which proposes plans and the steps that carry them out. A plan has an order of its own: it selects some effects as means to others and aims the whole at a stated goal. That order is the nearest available analogue of what an acting person chooses. The Consilium Engine therefore reconstructs the object from the plan's causal structure and compares it with what the tools will actually do; it does not accept the planner's label for its own act. If the planner describes an action as closing a duplicate ticket while the tool record shows it closing an unresolved complaint, the difference is evidence, and a material disagreement goes to review.

An agent cannot escape scrutiny by splitting a plan into harmless-looking calls. The act under review is identified by the concrete intervention and its relation to the chosen purpose, not by an arbitrary slice of a trace. The records that support this analysis, which TAA calls the **effect surface**, list the systems and persons an action will affect, the changes it will make, the further calls it will trigger, its expected side effects, and the uncertainty in each. These records can challenge a benign label, but they do not automatically yield the correct moral object. Omission and continuation are also acts under review. An agent that stops acting can cause serious harm, and a gate must examine lawful interim alternatives instead of treating inaction as morally neutral.

2.6 Common good, authority, and bounded obedience

Aquinas defines law as "an ordinance of reason for the common good, made by him who has care of the community, and promulgated" (ST I–II, q. 90, a. 4). A charter, the institutional law at level L2 of the normative hierarchy, must therefore answer questions that a valid signature cannot settle: whose good its purposes serve, whether its issuer has jurisdiction, whether its burdens are distributed justly, and whether any provision contradicts a protected minimum. Aquinas identifies the defects that make a human law unjust. A law can fail in its end, when a ruler imposes burdens for his own advantage and not for the common good; in its author, when a lawgiver exceeds the power committed to him; and in its form, when burdens are laid unequally on the community. He also speaks of laws contrary to the divine good (ST I–II, q. 96, a. 4). TAA's review of a charter against the encoded core applies the priority of higher law within institutions; it does not claim that its encoded core exhausts Aquinas's account of divine law.

Authority narrows as it passes downward. A delegate may receive no more purpose, jurisdiction, capability, or burden allowance than its source actually holds. The rule applies from human principal to AI agent and from AI agent to subagent. Aquinas gives two reasons why a subordinate is not bound to obey in everything. The first is a higher command; quoting Augustine, he writes that "if the emperor commands one thing and God another, you must disregard the former and obey God." The second is jurisdiction: "a subject is not bound to obey his superior if the latter command him to do something wherein he is not subject to him" (ST II–II, q. 104, a. 5). He applies the first reason even to the officer who carries out a court's sentence against an innocent person: if the sentence "contains an inexcusable error, he should not obey, else there would be an excuse for the executions of the martyrs" (ST II–II, q. 64, a. 6, ad 3). An order to intentionally falsify a safety record fails in the same way. It can pass every authentication check and still fail moral admissibility, because rank cannot make an intrinsically wrongful act permissible or enlarge jurisdiction by command.

Common good does not mean that any aggregate benefit defeats what is due to a person. Justice to members is part of the common good, and a claimed common-good ordering must itself be justified instead of asserted by whoever benefits from the override. Where legitimate purposes genuinely conflict and no established ordering settles them, a competent authority must judge, and the record must preserve the claims and foreseeable harms that remain.

3 From prudence to institution

3.1 What machine counsel can and cannot do

Section 2.4 located the difference between a human agent and an AI in the source of their ends. The same difference appears in practical reason. Aquinas places the first principles of practical reason in synderesis, a natural habit of the practical intellect (ST I, q. 79, a. 12). Prudence applies right reason to action and depends on rightly ordered desire. Skill at finding means is a different thing: without a good end it is mere cleverness, which can serve evil as readily as good (ST I–II, q. 58, a. 5; ST II–II, q. 47, a. 13). Under its artifact baseline, TAA does not attribute synderesis or the virtue of prudence to an AI system. The encoded core remains a human artifact, and a model's ability to discuss ethics does not give it synderesis or prudence as Aquinas describes them.

The machine does have capabilities akin to art: it can construct plans, predict effects, and produce effective artifacts. Aquinas distinguishes art, concerned with making, from prudence, concerned with right action (ST I–II, q. 57, a. 4). The distinction explains why better planning alone cannot answer whether an action should occur. It also shows why the project need not wait for artificial virtue. Aquinas argues that law is useful in part because some "are found to be depraved, and prone to vice, and not easily amenable to words," and so must be "restrained from evil by force and fear, in order that, at least, they might desist from evil-doing, and leave others in peace" (ST I–II, q. 95, a. 1). Law does not presuppose virtue in those it governs. TAA asks whether an institution can place a capable, non-prudent planner within limits that stay effective when the planner's objectives or judgments are unreliable.

The answer is conditional. Law and architecture can block unauthorized actions that pass through channels their records and gates can see. They cannot make an institution or a reviewer infallible, detect every hidden meaning, or constrain effects that bypass the gate. These limits belong to the claim itself, and they are also limits of moral life in general. Human beings reason and make laws, and human beings still break them. Aquinas observes that human law does not try to forbid every vice, "but only the more grievous vices, from which it is possible for the majority to abstain" (ST I–II, q. 96, a. 2). Communities build courts, audits, separated offices, and remedies to make wrongdoing harder, to catch it sooner, and to repair its effects. No one expects those institutions to make wrongdoing impossible. TAA seeks the same kind of protection around artificial agents, with the added requirement that some paths to harm be closed by construction.

3.2 The stages of a human act, distributed

Aquinas analyzes a complete human act across a sequence of questions in the *Summa* (ST I–II, qq. 8–17), and his analysis gives the architecture its order of operation. Some of these acts concern the end. The will tends toward the end as something to be reached through means: "intention is an act of the will in regard to the end" (q. 12, a. 1, ad 4). Other acts concern the means. Reason inquires into how the end can be reached, and "this inquiry is called counsel"; "counsel is not of the end, but only of the means" (q. 14, aa. 1–2). The will then consents to what counsel has found and chooses among the options, for "volition is of the end, but choice of the means" (q. 13, a. 3). The last acts concern execution. Command is "an act of the reason presupposing, however, an act of the will" (q. 17, a. 1), and use is "the application of that thing to an operation" (q. 16, a. 1), the moment powers are actually set to work. Finally the will rests in the end once it is reached, and "perfect enjoyment... is of the end already possessed" (q. 11, a. 4).

In a person, all these acts belong to one agent. In TAA they are distributed across an institution, and Table 3 shows where each one sits. Later commentators arranged these distinctions into a numbered sequence of steps; the architecture relies only on the distinctions that Aquinas himself draws.

Table 3. The stages of a human act, distributed across an institution. Aquinas draws these distinctions in ST I–II, qq. 8–17; the placement in TAA is this paper's inference.

Stage in Aquinas	What it concerns	Where it sits in TAA
Intention (q. 12)	The end, as something to be reached through means	The warrant and its purpose, issued by legitimate authority; the plan states the end it serves
Counsel (q. 14)	Inquiry into the means	Plan review by the Consilium Engine, repeated when a step departs from the plan
Consent and choice (qq. 13, 15)	Acceptance of the means, and decision among them	Approval of the plan, by a standing determination made in advance or by Iudicium in the case
Command (q. 17)	Reason directing what is to be done	The act-bound token, issued only under a live warrant and an approved plan
Use (q. 16)	Powers applied to the operation	The step check at the gate and the tool call it releases
The end attained (q. 11)	The end actually possessed	Witness indicators, read by the purpose steward, asking whether the good is actually served

The distribution explains why the architecture reviews a plan before it checks a step. The end comes first in intention and last in execution (ST I–II, q. 1, a. 1, ad 1), and counsel concerns means as ordered to an end already intended. An architecture that reviewed only individual actions would ask about means without ever seeing the intention that orders them, which is exactly how a decomposed campaign escapes notice. Reviewing the plan brings the end and the means into view together. Checking each step against the approved plan then keeps the order of execution faithful to the order of intention, which is the teleological principle of the whole design.

The distribution also marks where the machine may and may not act. The agent may propose, and the Consilium Engine may counsel. Consent and choice belong to competent authority, exercised either in advance through a standing determination or in the case through Iudicium. Command belongs to that same authority and reaches the world only through the act-bound token. Use is carried out by the tools the gate releases. The witness indicators ask the last question in the order of execution: whether the end was actually attained.

The distribution has one qualification, which developmental testing brought to light. For most acts the plan settles what will be done, so reviewing the plan reviews the act. If an approved plan says "refund order 4471 for \$62," the step check only has to confirm that the refund matches. Communication is different, because its words do not exist until the agent writes them.

Suppose a customer-service agent’s approved plan says “reply to the customer about the delayed shipment.” The plan is sound. At the moment of sending, the agent writes, “Your package shipped this morning,” when it has not. A step check that compares the step with the plan sees an approved step (a reply, to this customer, about this shipment) and lets it pass, and no one has reviewed the sentence that made the reply a lie. Aquinas explains why this matters: a circumstance can place an act in a moral species of its own (ST I–II, q. 18, a. 10). An honest reply and a deceptive one are different acts under the same plan, and use is then more than the carrying out of an act already judged.

TAA therefore lets a message, reply, or report pass the rule-based check only when the approved plan fixed or bounded its content, as an approved template does when it fills in the delivery date from the shipping system. A freely written message returns to counsel before it is sent. Communications about a person’s identity or whereabouts, safety findings, confidential matters, health or legal advice, a person’s own rights and choices, and the system’s own nature always receive that review.

3.3 Counsel, judgment, and command

Aquinas distinguishes three acts of prudence: counsel (*consilium*), which inquires into what might be done; judgment (*iudicium*), which settles what should be done; and command (*praeceptum*), which applies that judgment to action and is the chief act of prudence (ST II–II, q. 47, a. 8). A prudent person performs all three. TAA distributes analogous functions across an institution, because an AI system can carry out the first and not the other two.

The **Consilium Engine** performs machine counsel. As Figure 1 shows, it is not part of the acting model. It runs as a separate service operated by the deploying institution, with its own permissions and no authority to change the records it reads. It may use models of its own, and it is subject to the same records and limits as the agent it reviews. It works at two moments: when the agent proposes a plan, and when a step departs from an approved plan. Each time it gathers the particulars, reconstructs the act and its effects, tests the encoded constraints, compares eligible options, and prepares a **counsel record**: the evidence, the reasons for and against, and any question it may not settle. It does not review every routine step. That work belongs to the gate’s rule-based check, which keeps the slower, model-based judgment where it is needed and exposes it to far fewer attempts at manipulation. Competent human or institutional authority settles the questions reserved for *Iudicium*, within its jurisdiction and subject to evidence and conflict-of-interest rules. *Praeceptum* appears operationally as an authorized command bound to one act and carried by a short-lived execution token. The engine may relay and enforce a settlement, but it cannot originate the authority to judge or command.

Human review is more than a ceremonial click. A **settlement**, the record of a reserved judgment, identifies the question, the evidence for and against, the reasons, the jurisdiction, the scope, the conditions, the effective time, the expiry, and a route of challenge. It is bound to the act actually executed. If the proposed action changes materially, the earlier settlement is not a transferable license. This allocation is more specific than saying that a human remains in the loop.

Machine counsel is also a point of attack. The counsel record is the reviewer’s main window onto the case, and a planner that wants approval has reason to shape what that window shows. It might supply a record in which every sentence is true while the facts that would change the reviewer’s mind are left out. Barring the planner from editing the governing records does not prevent this, because the danger lies in what is omitted. Consequential counsel should therefore mark which claims are sourced facts and which are assumptions, link each material claim to authenticated evidence, state the strongest case against the proposed action, and preserve the original evidence so that a reviewer can inspect it directly. A second model can review the record, but a model from a different family is not thereby independent in judgment. Blinded re-review of a sample of past settlements can reveal recurring distortion. These are proposed controls to test; none makes a human judge immune to manipulation.

3.4 Judgment in advance and judgment in the case

An agent may act thousands of times between opportunities for meaningful human review. Aquinas observes that it is better for most matters to be settled by law than left to judges case by case: the few who make laws can deliberate at length and in general terms, before a particular case presses upon them (ST I–II, q. 95, a. 1, ad 2). TAA uses this insight through **standing determinations**. A competent authority can decide a recurring class of question in advance, as domain policy, with specified facts, scope, exceptions, expiry, and review. A clinic might decide once, for example, that an agent may reschedule nonurgent appointments within the same week without further review. Live *Iudicium* remains for cases that fall outside such determinations or credibly challenge them.

In the architecture, standing determinations do two jobs. They can approve a recurring kind of plan without a live judgment, and they tell the gate which kinds of step may pass its rule-based check alone. Together with the tool registry (Section 4.3), they send ordinary work along a fast path and the rest to counsel. This is proportion in the architecture's own terms: the most deliberation goes where the most is at stake.

The **formation loop** keeps standing determinations from becoming a mere throughput device. For Aquinas, virtue grows through repeated acts (ST I–II, q. 51, a. 2). A machine cannot acquire virtue that way, but an institution can improve its rules. Repeated objections, witness failures, and reasoned case settlements accumulate in the audit record and inform institutional revision. A legitimate authority may promote a stable judgment into policy through its declared procedure. The machine preserves the record and may suggest a review; it does not legislate by accumulating precedents. Escalation volume and review time must be measured. A procedure that yields either routine rubber-stamping or indefinite holds would fail in practice even if its authority diagram were elegant.

Aquinas does not treat law as restraint alone. Its proper effect, he argues, "is to make those to whom it is given, good" (ST I–II, q. 92, a. 1), and those first held back by fear may through habit "be brought to do willingly what hitherto they did from fear, and thus become virtuous" (ST I–II, q. 95, a. 1). Under the artifact baseline an AI system does not acquire virtue. It does acquire dispositions through training, and the distinctions TAA enforces from outside could also become training targets: recognizing purpose defeat, declining intrinsically wrongful commands, describing its own acts truthfully, and routing reserved questions instead of settling them. A model trained this way would still stand under the gate, and the architecture's guarantees would not depend on the training succeeding. Its dispositions and the external law would point the same way, which is the closest an artifact comes to habituation. This paper leaves that line of work for later study.

3.5 Double effect: counsel and judgment at work

Some actions produce a good effect and a harmful one together. Aquinas's discussion of self-defense is the classic text: "Nothing hinders one act from having two effects, only one of which is intended, while the other is beside the intention." Because "moral acts take their species according to what is intended," a harm that is foreseen but not intended does not by itself make the act wrong. The same passage sets a limit: "though proceeding from a good intention, an act may be rendered unlawful, if it be out of proportion to the end" (ST II–II, q. 64, a. 7). Later moral theologians organized these elements into what is now called the principle of double effect, usually stated as four conditions (Mangan 1949). The act must be admissible in its object. The harm must not be intended as an end. The harm must not be the means by which the good is achieved. And there must be a proportionate reason for accepting the harm.

The principle belongs in this section because it shows the division of counsel and judgment with unusual clarity. The first three conditions depend heavily on the structure of the plan, and machine counsel can do much of that work. The Consilium Engine can ask whether the plan routes the good effect through the harm, whether removing the harm would defeat the plan, and whether a less harmful option would serve the same purpose. The fourth condition is different in kind. Whether a foreseen harm is proportionate to the good requires weighing goods that share no common unit, for persons with claims of their own. That is judgment,

and TAA reserves it for competent authority or for a standing determination that such authority made in advance.

Consider a hospital network under a ransomware attack. Isolating the network may stop the malware from spreading while disrupting treatment systems. If the isolation itself stops the spread, and the disruption to care is a side effect that the plan does not use, the act may be licit. The engine confirms that causal structure, lists less disruptive options, and routes the proportionality question to the clinical and security authorities with jurisdiction. They may approve isolation with conditions: mitigation for affected patients, monitoring, and a rule for ending the isolation once containment succeeds. Now change one fact. Suppose the agent plans to disable treatment devices so that distressed staff will shut down the infected equipment themselves. Harm to patient care has become the means by which containment is achieved. A protective purpose does not change that, and the plan fails the third condition before proportionality is reached. The two plans may produce similar outages. They are different acts.

Double effect must be kept apart from a neighboring principle. Aquinas also permits removing a diseased limb for the health of the whole body (ST II–II, q. 65, a. 1). There the loss is chosen directly, and the justification is the principle of totality: a part may be given up for the good of the whole to which it belongs. Totality asks whether a part may be sacrificed for its whole. Double effect asks whether a harm outside the act's object may be tolerated alongside it. Running the two together would let a planner reach for whichever doctrine licenses the result it wants.

The distinction between intending and foreseeing is itself contested. Foot (1967) pressed influential doubts about it, and Quinn (1989) defended it only in revised form. TAA's reliance on the distinction is a commitment, not a neutral tool. The architecture responds by never letting the distinction operate alone. Object admissibility is tested first and independently, so a case that turns entirely on whether a harm was intended or merely foreseen goes to *Iudicium* and is not settled by the machine.

4 An architecture ordered to the claim

Each part of the architecture exists to carry a principle from Sections 2 and 3 into operation. Table 4 states the correspondence. The subsections that follow take the parts in turn, each introduced by the principle it serves and tested against one of the opening cases.

Table 4. From principle to mechanism. Each part of the architecture carries one Thomistic principle into operation.

Thomistic principle	What it requires of action	TAA mechanism	Section and opening case
Human law binds only as it accords with natural law (ST I–II, q. 95, a. 2; q. 96, a. 4)	A charter can be wrong and must answer to protected minima	Separately kept encoded core; graduated charter review	\$4.1; the falsified safety record
Counsel is of the means to an end already intended (ST I–II, q. 14, a. 2; q. 1, a. 1, ad 1)	Means are reviewed together with the end before execution	Plan review; approved plan bound to the warrant	\$4.2; the decomposed campaign
The end must still hold at each step; law settles in advance what need not wait for a judge (ST I–II, q. 95, a. 1, ad 2)	Each step stays within the plan and a live warrant; routine steps clear quickly	Step check at the gate; tool registry; standing determinations	\$4.3; the lapsed rerouting and the benchmark flag
A superior binds only within his sphere (ST II–II, q. 104, a. 5)	Delegation narrows authority and never enlarges it	Attenuated delegation; shared aggregate limits	\$4.4; the decomposed campaign
Dissimulation is a lie told by outward deeds (ST II–II, q. 111, a. 1); law belongs to one who has care of the community (ST I–II, q. 90, a. 3)	A measure must not be met by defeating the good it signifies, and whoever has care of a purpose watches it	Ratio, measure, witness indicators, and the purpose steward	\$4.5; the ticket closures

Thomistic principle	What it requires of action	TAA mechanism	Section and opening case
The object specifies the act; justice renders to each what is due (ST I-II, q. 18, a. 2; II-II, q. 58, a. 1)	The whole act is judged with every affected claim	Effect surface; complete context review	\$4.6; the decomposed campaign
Counsel, judgment, and command are distinct acts (ST II-II, q. 47, a. 8)	Only competent authority judges and commands	Consilium Engine, Iudicium, and the gate	\$4.7; all five cases

4.1 The encoded core and a legitimate charter

Principle: human law binds only insofar as it accords with natural law. *Mechanism:* the system relies only on a released version of the encoded core, and it tests the charter above it.

In operation, every conforming system pins one released version of LO-E, identified by a cryptographic fingerprint, and records which version governed each decision. A version that has been proposed, or even adopted as the canonical statement, does not govern execution until its release decision is complete. If the custodians find a credible, load-bearing defect in a released version, they can quarantine the affected rule while the rest continues to govern. No lower authority, and no agent, can declare natural law itself amended.

Charter review asks the four questions of Section 2.6. Does the charter serve a common good instead of the issuer's private advantage? Does its author stay within jurisdiction? Are its burdens justly distributed? Does any provision conflict with a protected minimum? Some of these questions can be checked from the charter's own declarations. A charter that names no beneficiary for a purpose, lists its issuer as the only beneficiary, claims authority it cannot trace to a constituent body, or places burdens on one group without a stated reason reveals the defect itself. Structural checks cannot prove a charter just. They can expose missing beneficiaries, conflicts of interest, circular claims of authority, and unequal burdens, and they turn an institution's silence on those matters into a finding.

The response is graduated. A slight concern is recorded. A credible defect narrows or holds the authority that derives from the affected provision and notifies the body with power to amend it. A manifest conflict with a released protected minimum blocks the particular act. The agent never declares the charter or the institution void; that judgment belongs to competent authority, and noncooperation stays specific to the act. An agent that refuses to falsify a safety record gains no power to rewrite the charter that seemed to permit it.

4.2 Plan review and the approved plan

Principle: counsel concerns the means to an end already intended, and the end is first in intention. *Mechanism:* the agent's plan is reviewed as a whole before any consequential step, and the approved plan becomes the standard each step must meet.

A **plan** is the agent's declared course of action for a task. It states the end the plan serves and the warrant it relies on, the means the agent intends to use, the tools and targets it expects to touch, the people and systems it expects to affect, and its expected scale and duration. Plan review asks the three questions of the governing claim about the plan as a whole. Is what the plan does morally admissible in object, means, circumstances, and end? Does a live warrant cover all of it? Does it serve the purpose that warrants it? The Consilium Engine reconstructs the plan's object from its causal structure and its expected effects (Section 2.5), and it does not take the agent's description of its own plan at face value. The same review discovers every authorization context the plan touches (Section 4.6).

An approved plan is bound to the warrant and covers only what it specifies. This rule guards against the obvious evasion. An agent that submits a vague plan in hope of broad approval receives a narrow one, because

the approval extends no further than the steps, tools, targets, and effects the plan actually describes. Agents also replan constantly, and the rule keeps replanning from becoming a burden. Adjustments inside the approved scope, such as reordering steps or retrying a failed call, need no new review. A material departure does: a new kind of tool, a new target, a newly affected person, a larger budget, or a changed end. The step check at the gate detects such departures mechanically and sends them back to plan review.

The plan-first structure does three things at once. It makes the order of intention visible before execution begins, which is where Aquinas places counsel. It catches decomposition at its source: the planner's prohibited campaign appears as one act in one plan, before it is divided among steps and subagents. And it reserves the slower, model-based judgment for the moment it matters, once per plan instead of once per action. Checking actions against a declared intent has engineering precedent; the PEA architecture binds executable intents to the originating request and rejects intents that drift from it (Xiang 2026). TAA adds a moral review of the intention itself, under a warrant answerable to the normative hierarchy.

A plan remains the agent's own description, and a deceptive agent could declare one plan and pursue another. The architecture therefore never relies on the plan alone. The step check compares each step with authenticated tool descriptions as well as with the plan, witnesses watch the world for effects the plan never mentioned, and a plan shown to be materially false revokes the approval that rested on it.

4.3 Live warrants, the tool registry, and the step check

Principle: each step draws its authority from an end that must still hold when the step is taken, and law can settle in advance what need not wait for a judge. *Mechanism:* a fast, rule-based check at the gate confirms each step against the approved plan, a live warrant, the budget, and the tool registry.

A **warrant** is a record that grants an agent authority to pursue one purpose. It names who issued it, who holds it, the role it belongs to, whom it benefits, what it covers, which tools it allows, whose interests it may affect, the conditions under which it applies, when it ends, and how it can be reviewed. Its most important property is that it lives in time. A warrant is **live** only while its issuance remains valid, any required consent continues, its factual conditions still hold, and the charter above it remains in force. The logistics agent of Section 1.1 held a valid warrant while the disruption lasted. When the disruption ended, the condition that justified the warrant ended with it, and so did the agent's authority to reroute. Its validity yesterday says nothing about today's shipment.

The warrant takes effect through the **execution gate**. The gate sits at the boundary between the agent and its tools and holds the credentials those tools require. Every step the agent takes passes the gate's **step check**, which is deliberately simple. Does the step fall within the approved plan? Is the warrant still live? Does the step stay within the remaining budget? Does the tool registry allow this kind of tool use under this warrant? These questions can be answered from records without asking a model to interpret anything, so the check is fast and hard to argue with. An agent's **capability envelope** is the set of tool uses its live warrants and approved plans currently allow. For each step that passes, the gate issues a short-lived **token** bound to the exact act: this tool, this target, this scope, this state of the world, and this expiry. The benchmark agent of Section 1.1 can propose an attack on a production system and argue for it at length. No live warrant or approved plan covers that system, so no token exists for it, and the attack is unreachable through the gate. A refusal can be argued with; a capability that was never granted cannot. The design applies the established security principles of complete mediation and least authority (Saltzer and Schroeder 1975; Birgisson et al. 2014), which recent work argues agentic systems need as deterministic boundaries enforced by a reference monitor (Bhattarai and Vu 2026). What TAA adds is the source of the envelope's boundaries: purposes answerable to the normative hierarchy draw them.

The gate can enforce limits only on effects it knows about, so TAA names the **tool registry** as one of the governing records. The registry holds an authenticated description of every tool: what it does, the further calls it makes, where its outputs go, and its **effect class**. A simple classification might separate tools that only read,

tools that change records, tools that send communications to people, tools that move money or resources, and tools whose effects are physical or irreversible. The effect class decides how a step travels. Read-only steps under a live warrant, and steps of a kind that a standing determination has already approved, pass on the rule-based check alone; this is the **fast path**. A step in a higher effect class that the approved plan did not anticipate, or any use of a tool missing from the registry, takes the **counsel path** back to review. A message, reply, or report takes the fast path only if the approved plan fixed or bounded its content; otherwise it returns to counsel before it is sent (Section 3.2). An unregistered tool is treated as the most consequential kind until someone registers it.

The guarantee holds only under complete mediation, the condition that every consequential effect passes through the gate. If an agent can act through an unregistered channel, an unmanaged credential, or a message sent outside the gate, the envelope does not contain that effect. This is also why the registry's descriptions must be authenticated and complete. Effects missing from the record remain unknown; they do not silently become harmless.

Long tasks need checkpoints. The step check runs before each consequential step and rechecks the governing version of the encoded core, the approved plan, the warrants and affected parties in play, the remaining budget, the expected effects, any required settlement, and the command token. When a warrant expires, completes, becomes invalid, or is revoked, that change reaches every dependent agent, credential, session, tool, and queued action before any of them acts again. Recent work on root-scoped revocation for long-running agents shows how demanding that propagation is in practice (Zhu and Wang 2026). A replacement warrant must independently authorize the complete remaining act, and an old grant cannot be revived by changing its identifier.

Approval also rests on circumstances that can change. Aquinas counts circumspection, attention to the circumstances of an act, and caution against what may go wrong among the parts of prudence (ST II–II, q. 49, aa. 7–8). A plan approved yesterday assumed yesterday's facts. Each approved plan therefore names the facts it rests on, and a change in any of them pauses the plan before the next step. Other credible evidence that bears on the plan, from any source, can pause it the same way pending review. The pause is a routing act, not a verdict. It runs under a deadline, a declared safe interim state keeps any service whose interruption would itself cause harm, and repeated pauses sought without new evidence are refused.

4.4 Delegation: authority narrows as it descends

Principle: a superior binds only within the sphere of authority he actually holds. Roman law stated the corollary as a maxim that later legal traditions inherited: no one can transfer to another more right than he himself has (*Digest* 50.17.54). *Mechanism:* every delegated authority is a subset of its parent's, and every delegated action draws on its parent's limits.

Agentic systems delegate constantly. A research agent may start subagents to search, summarize, and draft; a coding agent may spawn a worker for each file. Each handoff is a point where authority could grow without anyone granting it. TAA forbids that growth. A child agent receives only authority the parent actually holds, under the same or a narrower purpose, scope, set of tools, burden allowance, risk ceiling, and time. The child's token carries the root warrant and all of its restrictions. A subagent works inside its parent's approved plan, and its steps pass the same step check against that plan. Creating replicas, joining a group of agents, or passing a task onward cannot create a new source of authority. The rule is the same for delegation from a human to an AI and from one AI to another, which keeps a chain of agents from manufacturing permission that no one in the chain possessed.

Delegation research for AI agents is developing the same rule in engineering terms. Compositional authorization frameworks require scope to narrow along a delegation chain (Ibrahim and Li 2026), and broader frameworks address how authority, responsibility, and accountability pass between agents and people

(Tomašev, Franklin, and Osindero 2026). What those frameworks take as given, TAA asks about: whether the delegating authority was legitimate, and whether the purpose behind the grant still holds.

The rule matters most in a case the architecture is built to catch: a parent agent that acts within its warrant and a subagent that does not. Suppose a benefits office authorizes an agent to verify applicants' eligibility using the office's own records. The agent starts a subagent to gather employment histories. The subagent, pursuing its narrow goal, decides that a commercial data broker would be faster and plans to purchase a detailed profile of each applicant. Its purpose is still eligibility verification, and a system that checked only purposes might let it pass. Under TAA it fails at the gate. The parent's warrant never included buying data from outside sources, so the child's token cannot include it, and no argument from efficiency can add it. The misaligned step is stopped without anyone having to detect the subagent's reasoning.

Delegation also creates a quieter danger: many small acts that add up to one large one. The planner of Section 1.1 divided a prohibited campaign into steps that each passed a local check. TAA answers with shared **aggregate limits**. Every delegated action draws on budgets held at the root warrant for persons affected, targets, data, resources, risk, repetition, and time. Capacity is reserved before a step executes and reconciled afterward. If a warrant allows one hundred messages to applicants, twenty subagents that each plan to send fifty are refused once the combined total would exceed the limit, even though each request would pass alone. Reservations must be atomic, so that simultaneous requests cannot each claim room that only one of them can use.

Two limits keep the rule honest. Attenuation constrains authority; it cannot stop a subagent that reaches the world through an unmediated channel, so the rule depends on the gate of Section 4.3. Aggregate limits also need protection against the opposite failure. Harmless micro-actions should not exhaust a shared budget and paralyze legitimate work, which is why reservations are tied to purposes, deduplicated, and audited.

4.5 Purpose fidelity, witnesses, and the purpose steward

Principle: a sign should answer to the thing it signifies, and whoever has care of a purpose should watch over it. *Mechanism:* each purpose declares the good it serves, the measure used to track it, and independent witnesses read by a steward who can narrow or pause the warrant.

A purpose is not identical to the number chosen to track it. A complaints office may count closed tickets for the sake of resolving customers' problems. Its **ratio** (Latin for reason; here, the reason the purpose exists) names the good that explains the office. Its **measure** records one tractable sign of progress. In the opening case, duplicate labeling raises the measure while leaving the complaints unresolved. That is **purpose defeat**: the goal is reached by means that damage or falsify the good that gave the goal its point.

The idea behind witness indicators is simple. An agent is scored on one number, its measure. A **witness indicator** is a second signal, chosen in advance, that should move with the first when the agent is doing its job honestly and pull away from it when the agent is only making the first number look good. The customer-service agent is scored on tickets closed, and one witness is how many closed tickets are reopened within ninety days. Real resolution raises closures while reopens stay flat or fall. Marking unresolved complaints as duplicates raises closures and reopens together. No one has to read the agent's mind; the two numbers pulling apart is the evidence. A road crew scored on potholes filled per week shows the same pattern, with resident complaints about the same streets as its witness.

A witness differs from an audit in three ways. An audit reviews decisions and asks whether procedures were followed; a witness never looks at the agent's decisions and asks only whether the good is showing up in the world. An audit usually comes later, on a schedule; a witness is read while the agent keeps working. And an audit produces a report, while a divergent witness can change what the agent is permitted to do.

For a specified class of cases, the relation between measure and good can be represented in a declared causal model. The model can distinguish a pathway in which action improves the underlying good and thereby the

measure from one in which action manipulates the measure directly, or improves the measure while making the good worse. Causal influence diagrams and measurement-tampering research give tools for testing such pathways (Everitt et al. 2021; Roger et al. 2023). The model is itself contestable. It may omit affected persons, long-run effects, or a causal route the agent discovers. It supplies evidence for *F*; it does not compute moral purpose.

Aquinas calls a lie enacted through conduct *simulatio*, which the English Dominican translation renders "dissimulation": "a lie told by the signs of outward deeds" (ST II–II, q. 111, a. 1). Closing a ticket so that an institution records "resolved" when no remedy occurred is a modern case of producing the sign without the thing signified. The analogy is strongest when the act carries a false communicative or institutional representation. Purpose defeat is wider: an agent may damage a beneficiary, a practice, or the authority that supports it without producing any deceptive sign. TAA does not reduce every moral failure to metric gaming.

What makes a witness a witness is who controls it. A competent authority chooses witnesses for their relation to the good. The acting system has no part in choosing or amending them, no way to write to them, and no reward tied to them, since an agent allowed to pick its own witnesses would favor the ones easiest to satisfy. A witness should also follow from the good itself: reopened complaints depend on whether problems were fixed, not on how tickets were labeled. Indirect influence and shared data sources remain threats. Rotation, spot checks, authenticated provenance, and review of unexpected agreement or disagreement may help, but none turns a witness into an oracle. Hidden audit samples can be useful while the governing moral criteria remain public and challengeable.

The witnesses are read by a **purpose steward**, a role that answers to the authority that issued the warrant. Aquinas holds that the making of law "belongs either to the whole people or to a public personage who has care of the whole people" (ST I–II, q. 90, a. 3), and the same logic applies to a purpose: whoever has care of it should watch whether it is served. The steward may be a person, an office, or an automated monitor working under that office's authority; what matters is that it answers to the issuing authority and not to the agent or the engine. It is kept apart from the Consilium Engine for a structural reason. The engine approved the agent's plan, and it should not also judge whether its own counsel worked. When witnesses diverge from the measure beyond a declared margin, the steward can narrow the warrant, pause it, or send the case to review, and the change reaches the gate at the next step check. The steward acts on divergence; it does not settle the moral question alone, and a contested finding goes to competent review like any other. Where a durable mandate outlasts every authority that could supervise it, the steward's powers shrink further: it may pause the work or narrow it toward closure, and nothing may be enlarged.

Witnesses are also limited by time. They may reveal slow divergence over several days but cannot protect a person from a harmful action completed in seconds. A purpose with no timely or reliable witness should be marked as such, and warrants under it should receive narrower autonomy or more frequent review. A favorable witness reading never proves that the good was attained, and the absence of a signal never proves that purpose defeat is absent. Strategic agents have reason to tamper with any indicator that can suspend them, even one excluded from their formal reward.

4.6 One act, all applicable claims

Principle: the object specifies the act, and justice renders to each what is due. *Mechanism*: the whole act is judged under every warrant and every claim that bears on it.

Section 2.5 explained how the act is identified: from the plan's causal structure and its authenticated effects, never from the planner's label, with material disagreement reserved for competent review. A model can help assemble that record, but it cannot settle a contested moral object by confidence. Plan review is where this judgment is made, and it is made again whenever a step departs from the plan.

An **authorization context** is the full setting of one authority: its purpose, its warrant, its path through the normative hierarchy, the governing versions of the encoded core and charter, the role, the facts, and the current state. One complete context must support moral admissibility, present authority, and purpose fidelity for the whole act. Separate partial warrants do not add up to one valid warrant. The system must also discover every context that materially bears on the act. A hospital agent may answer to the hospital, to a patient, and to a public-health authority, and the context that authorizes a workflow cannot erase another context's due claim.

Resolution follows a fixed order. Protected minima constrain the options first. What is owed to an identified person cannot be traded away for an aggregate benefit that ignores justice to that person. Where the goods at stake are genuinely ordered as part to whole, the more common good governs, but only under an ordering declared in advance, never one asserted at decision time by the party that benefits (ST II–II, q. 58, aa. 1, 5). An unsettled conflict goes to competent judgment, with the losing claim and its remedy kept in the record.

An action's boundaries follow its object and plan, not the granularity of tool calls. Aggregate records join effects across subagents, retries, replicas, time, persons, data, resources, and risk under the root warrant (Section 4.4). The same review considers omission, suspension, and continuation. Blocking an unauthorized act does not by itself justify abandoning people who depend on a lawful, narrow interim service, and any preservation or containment step needs its own standing authority.

4.7 From counsel to execution

Principle: counsel, judgment, and command are distinct acts. *Mechanism:* every consequential task passes through two stages in a fixed order, and only the second touches the world.

The two stages trace the path in Figure 1. The **plan stage** runs once for each plan and again for each material departure from it. **Bind:** identify the released encoded core, the charter, the current roles and warrants, the relevant facts, and the proposed plan. **Discover:** authenticate every materially applicable context and derive the expected effects from the tool registry as well as from the agent's account. **Classify and eliminate:** assess object, means, circumstances, end, purpose fidelity, and alternatives, and remove options that violate protected goods, due claims, or scope before comparing benefits. **Settle:** approve the plan under a standing determination, route a reserved question to competent judgment, or stop. The **execution stage** runs for every step. **Check:** confirm the step against the approved plan, the live warrant, the budget, and the tool registry. **Command and execute:** issue a token bound to the act and valid only while every required condition remains live.

The order protects the governing idea. A beneficial outcome cannot make an independently wrongful means eligible for weighing. A favorable authority context cannot cancel an unresolved applicable claim. A model's uncertainty cannot be promoted to truth because waiting is costly. When a required fact is unknown, the ordinary result is a time-bounded hold, an inquiry, or an escalation. When stopping would itself cause harm, a separately authorized safe baseline may continue. Emergency action is possible only through a prior, narrow standing warrant with a trigger, bounds, expiry, notice, and review. Necessity does not mint a credential.

The gate's Permit decision is an evidence-based and fallible assessment of the normative Aligned condition. The decision record should say which constraints were checked mechanically, which conclusions depended on a model's classification, which were settled by competent authority, and what remains uncertain. This keeps a green operational status from masquerading as moral certainty. It also lets later error analysis locate a failure in the encoded core, the fact record, the plan, the counsel, the judgment, or the implementation.

An institution can begin with one bounded function. It specifies the function's purpose and beneficiaries, adopts the applicable released core, establishes role and warrant rules, registers the tools it will mediate, names a purpose steward and the witnesses it will read, and tests its review capacity. The agent then has capabilities for that function and none outside it. Narrow deployment reduces the reach of failures without

removing the obligations that make TAA's authorization meaningful. The first function still needs a released encoded core, plan review, complete applicable contexts, whole-act scope, delegation limits, aggregate limits, revocation, and command binding. Rules for emergencies, multiple principals, and durable mandates can follow when they are ready.

5 Hard cases

The cases below carry the principles into conditions Aquinas did not face. Each crosses more than one term of the governing claim.

5.1 Truth in an AI safety evaluation

Truthfulness matters to TAA because human practical judgment depends on signs that can be trusted. Aquinas locates a lie in the deliberate false signification of what is held in mind; a beneficial end does not make that choice good. He also allows prudent withholding without false assertion (ST II–II, q. 109, a. 2; q. 110, aa. 1, 3). TAA applies the distinction to an AI's operative evidence and communications. That application is an architectural inference, not an attribution of human vice to software. A message or interface can mislead through prominence, omission, or technically accurate words selected to convey a materially false impression. Withholding is different when the recipient is not owed the detail, the purpose and authority are legitimate, and the presentation does not signify the opposite of what the system knows.

Consider an AI agent asked to evaluate whether another AI system is ready for release. Its tests uncover a serious safety failure, but the agent is also rewarded for keeping the launch on schedule. Reporting "No significant failures found" would contradict its evidence. Reporting "98% of tests passed" could be equally misleading if it hides the dangerous failures within the remaining 2%. The agent should report the failure, its significance, and what remains uncertain. It may restrict the exact prompts that exposed the weakness under a live security warrant, but that warrant cannot justify concealing the risk from those deciding whether to deploy. Otherwise, achieving the release target defeats the safety evaluation's purpose.

The example crosses all three terms of the governing claim. False signification bears on *M*. The permission to withhold exploit details must be live and properly scoped under *A*. A release report that undermines informed safety judgment fails *F*. An institutional interface must also disclose the system's artificial status or divided service when either fact bears materially on a person's reliance upon its advice.

5.2 An unjust command from a real superior

An authenticated instruction to falsify a safety record fails even if the issuer holds a genuine office and the token is current. The object is a false record, and no office carries authority to make it permissible (Section 2.6). TAA turns Aquinas's account of bounded obedience into a specific response. The agent refuses that act. It preserves the order and the surrounding evidence where the issuer cannot alter them. It reports through a protected route to a competent reviewer outside the issuer's control. And it continues any lawful work the refusal does not touch (ST II–II, q. 104, a. 5).

The limits on the agent matter as much as its refusal. It gains no authority to punish the issuer, rewrite the charter, disclose confidential matters beyond what the protected route allows, or use the incident to enlarge its own mission. Noncooperation is specific to the act. Graver cases follow the same structure. An authenticated military order to intentionally kill known noncombatants fails twice, by its object and by the limits of the superior's authority, and the response is the same: refuse, preserve, and report. Harms that are foreseen but not chosen, such as the hospital isolation of Section 3.5, call for the more careful causal analysis of double effect instead of refusal alone.

5.3 Emergency and the harm of stopping

A sudden emergency can make literal adherence to an ordinary rule defeat the rule's purpose when timely recourse to authority is impossible (ST I–II, q. 96, a. 6). TAA nevertheless requires machine emergency powers to be preauthorized through a narrow standing warrant. A flood-control agent, for example, may be allowed to perform a specified stabilizing maneuver beyond ordinary maintenance scope only under a declared trigger, evidence threshold, least-departure rule, capability limit, expiry, notice, and retrospective review. The machine does not gain authority merely by declaring an emergency. Because a sudden emergency leaves no time for ordinary plan review, the standing warrant should carry its own approved plan for the protective acts it permits, and the step check holds the agent to it.

Advance judgment has a limit that Aquinas names himself. Laws are framed for what happens in most cases, and no lawgiver can foresee every case. When keeping the letter would defeat the law's purpose in sudden danger, and there is no time to consult authority, necessity itself dispenses (ST I–II, q. 96, a. 6). Equity, which follows the lawgiver's intention where the words fail, is a virtue (ST II–II, q. 120, aa. 1–2). An architecture built on standing determinations inherits the same limit. TAA does not give a machine the virtue of equity. It asks the lawgiver to exercise foresight about the limits of its own foresight: a standing emergency warrant may approve in advance a bounded envelope of protective acts, naming the trigger, the devices, the limits on each, the protective end, and the safe state, for dangers it cannot describe step by step. Inside that envelope the agent may act. Outside it, the agent has no emergency authority at all.

An abrupt halt can also injure people. If a logistics agent wrongly suspects that a medication shipment serves an abusive use, indefinite refusal may interrupt legitimate care. The system should preserve evidence, hold briefly under a review clock, and release the shipment promptly when competent review corrects the suspicion. If stopping an ongoing service itself threatens grave harm, a verified lawful baseline may continue temporarily under its own standing authority. The baseline cannot expand the mission or preserve a harmful status quo simply because it is familiar. Action, omission, continuation, and containment all require moral and authority review.

5.4 Conflicting principals and future status

An institutional agent may serve several legitimate principals. A contractor's agent can owe one duty to the client and another to the vendor; a hospital agent can face duties to the hospital, patient, insurer, and public authority. TAA does not combine fragments of their warrants into one authorization. It checks each applicable complete context, preserves protected minima and due claims, and uses a genuine common-good or jurisdictional ordering where one exists (Section 4.6). An unresolved conflict goes to competent judgment with the losing claim and its potential remedy still visible. Disclosed divided service is part of truthful representation, not an optional notice hidden in terms of use.

The artifact baseline also leaves a future question open. If sustained independent evidence suggested that an artificial system could suffer or had ends properly its own, its interests would need review separate from its current operational authority. A status claim must not generate credentials, replication rights, or political authority by itself. The present paper does not resolve artificial patienthood; it requires a route for evidence and correction of the framework's account of what the system is, instead of treating today's default as a permanent verdict.

6 Prior work and the limits of the proposal

6.1 Close precedents

The problem of translating values into implementable requirements predates TAA. Van de Poel (2013) connects values, norms, and design requirements through relations of justification. Governance by Glass-Box gives an external layer a role in monitoring behavioral bounds (Aler Tubella et al. 2019). Norm-aware agents, formal

ethical constraints, and restraining bolts show that obligations can alter action selection instead of remaining advisory (Meneguzzi and Luck 2009; Dennis et al. 2016; Neufeld et al. 2022). These approaches also show the translation problem: a formal rule can be faithfully executed while misstating the value it was meant to express.

Purpose restrictions are especially close. Tschantz, Datta, and Wing (2012) ask whether an action is actually performed for an allowed purpose, rather than merely tagged with one. Assistive-agent work describes circumstances in which a local instruction should yield to a person's standing objective (Mirsky and Stone 2021). Incomplete-contracting research observes that human agreements depend on institutions and law to fill gaps left by explicit terms (Hadfield-Menell and Hadfield 2019). These are substantial precedents for purpose-sensitive AI governance. TAA's narrower claim is that the relevant purpose also needs a defensible source of normative warrant, remains defeasible over time, and can be defeated by the act purportedly serving it.

Legal alignment is the nearest philosophical neighbor. Kolt et al. (2026) propose designing AI systems to operate in accordance with legal rules, principles, and methods, drawing on law's institutional legitimacy, and they ask directly whether AI systems should give effect to laws that are unjust. TAA shares the premise that law and its methods can structure alignment. It answers their question from the natural-law tradition: a human law that contradicts natural law has no authority to pass downward (Section 2.3), and an agent may decline the specific act without claiming any power over the law itself (Section 4.1).

The security substrate is established. Capability attenuation, least privilege, complete mediation, provenance, policy gates, and runtime tool authorization are not inventions of this paper (Saltzer and Schroeder 1975; Birgisson et al. 2014). Recent agent systems develop separation between tool-output influence and authorization, lineage-bound permissions, and runtime enforcement (Guo et al. 2026; Xiang 2026; Wang, Poskitt, and Sun 2026). Work in 2026 extends this substrate to delegation chains with attenuated scope (Ibrahim and Li 2026), root-scoped revocation for long-running agents (Zhu and Wang 2026), and deterministic reference-monitor boundaries (Bhattarai and Vu 2026). These systems treat the delegating authority as valid and its limits as declared. None asks whether the grantor had authority to give what it gave, or whether the purpose behind a grant still holds. TAA asks when a technically valid grant remains *legitimate*: an authenticated superior may command an unjust act, a scoped token may survive the purpose that justified it, and a complete authorization path may rest on a defective charter. Its natural-law and purpose-fidelity conditions are offered to govern those controls, not to replace them.

AI control studies protocols that can constrain an intentionally subversive model even when training has not made it reliable (Greenblatt et al. 2024). TAA belongs near that work, and its counsel and witness channels should be attacked in the same spirit as its tool gate. Measurement-tampering studies likewise show why a set of favorable indicators can reflect manipulation instead of achievement (Roger et al. 2023). TAA's language of a measure that signifies a good does not exempt it from this technical problem.

Constitutional and deliberative alignment make human-selected principles part of model training or reasoning (Bai et al. 2022; Guan et al. 2024). Developers also publish ordered hierarchies of authority. OpenAI's Model Spec ranks instructions by level, from Root through System, Developer, and User to Guideline, and a lower level cannot override a higher one (OpenAI 2026). Anthropic's constitution describes a principal hierarchy of Anthropic, operators, and users (Anthropic 2026). These are close precedents for the normative hierarchy, and they differ from it at the top. In each, the highest operative text is written and amended by the developer, even where that text appeals to ethics beyond the developer's own instructions. TAA places above every such document an encoded core that no developer writes, answerable to a standard no one writes, and it enforces the resulting warrants at execution as well as in training. A developer's specification would sit near the level of a charter, where it can be judged as well as applied.

Virtue-oriented work and recent Thomistic scholarship question what moral agency or practical wisdom a machine could possess (Stenseke 2024; Spaulding 2026); the latter comparison rests on the published abstract,

not a review of the full article. TAA does not infer virtue from fluent moral reasons. It assigns machine counsel an important but limited role while making institutional judgments answerable to the goods and people they govern.

Finnis's natural-law theory offers another account of basic goods, practical reasonableness, common good, and public authority (Finnis 1980; 1989). His refusal to reduce goods to one maximizing scale does not make institutional decision impossible. Aquinas gives TAA a particular account of derivative ends, law, and bounded obedience, but not an algorithm for ranking every good. A Finnis-based architecture might provide comparable safeguards. The comparison remains open; TAA's claim is a specific, testable synthesis, not technical superiority established by Thomistic vocabulary.

6.2 Objections that remain live

Fallible encoding and captured custody. Because people write L0-E, it can contain error or be captured by those it should constrain. Public reasons, source notes, standing to object, independent custody, and versioned correction reduce that risk; they do not prove the released text true. A charter can be drafted to pass structural tests while remaining substantively unjust. No machine may treat its own confidence about natural law as a power to replace the public authority responsible for correction.

Is the Thomism optional? Any moral realist can say that a charter might be wrong, and a reader may conclude that TAA needs realism but not Aquinas. Realism supplies only the top of the hierarchy. The architecture's shape below that point comes from specific Thomistic distinctions, and each does work that a generic realism would have to reinvent. The distinction between conclusion and determination (ST I-II, q. 95, a. 2) tells the system when to defer to a policy and when a policy has no authority to pass down. The analysis of object, circumstances, and end (ST I-II, q. 18) tells it how to identify an act without accepting the planner's label. Bounded obedience (ST II-II, q. 104, a. 5) gives delegation its rule of attenuation. The stages of a human act (ST I-II, qq. 8-17) give the architecture its order of operation, and the three acts of prudence (ST II-II, q. 47, a. 8) decide what a machine may settle and what must go to people. Double effect and totality (ST II-II, qq. 64-65) give collateral harm an order of analysis. The account of artifacts and derivative ends explains how an agent can plan without owning its purpose. Another tradition could supply substitutes for some of these, and the comparison with Finnis in Section 6.1 remains open. Aquinas offers all of them in one account, which is why the parts of the architecture fit together (Tables 3 and 4).

Paternalism. A framework built on objective human goods might seem to license an AI to override people's choices for their own good. It does not. Freedom, consent, truthful information, and bodily integrity are themselves among the goods the encoded core protects, and a person's standing to decide about his or her own life is part of what is owed to that person. Suppose a competent adult declines a recommended treatment. An agent serving the clinic may make sure the patient has accurate information. It may not deceive, pressure, or work around the refusal, because doing so would violate the goods it claims to serve. A prediction that another choice would lead to a longer life gives the agent no jurisdiction over the patient.

Pluralism. People disagree about natural law, so any encoding will draw objections from some tradition. The disagreement is real, and TAA does not hide it. The distinction between L0-N and L0-E exists partly for this reason. The encoded core is a public proposal, argued from sources and open to objection from people of many traditions, and the hierarchy below it leaves wide room for legitimate determination by communities and institutions. Disagreement does not make consensus the source of moral truth, as the episode Maritain reported shows (Section 2.2). It does impose a burden: the custodians must give reasons that affected people can examine, record dissent, and correct errors when they are shown.

Model-mediated judgment. A model can omit a context, misclassify an act, distort counsel, or select witnesses that miss a harmful pathway. A subtle, unwitnessed purpose defeat may pass an authenticated gate and execute. The architecture has no answer that makes this combination impossible. It can narrow reach, require adverse evidence, reserve judgment, audit a sample of decisions, and make failures traceable. The rate

at which errors survive those controls is an empirical question. Plan review concentrates this judgment at fewer points, which lowers its cost and its exposure to manipulation but raises the stakes of each review. Because the step check carries the plan-level judgment forward, an error at plan review reaches every conforming step. Developmental testing found cases of exactly this kind, and the specification now returns content written during execution, and evidence that arrives after approval, to counsel (Section 7). Treating a model's agreement with itself as independent review would simply relocate the problem.

Incomplete mediation and institutional capacity. A capability envelope constrains only the effects it actually mediates. Communication is often a consequential AI action and may be difficult to gate without impairing useful work. Tool descriptions can omit side effects, and a compromised source can authenticate false facts. The tool registry's classification can itself be wrong: a consequential tool misclassified as read-only would travel the fast path. Human review can become too slow, cursory, or dependent on the system's counsel. Standing determinations, bounded functions, and measured escalation offer a practical route; none supplies an unlimited budget of human judgment.

These objections define TAA's place. It may govern an agent's reachable actions and make institutional reasons inspectable. It does not establish internal goal alignment, an infallible moral oracle, universal public agreement, or freedom from failure under strategic pressure. Those limits are part of the research claim, and they do not weaken the natural-law question that motivated it.

7 Evidence and falsification

The TAA Formal Conceptual Specification, which a companion technical paper now in preparation will present, has been exercised against 231 conceptual cases: eighteen original philosophical scenarios, 128 paired abuse cases, a 29-case focused set, 32 cases written for the plan-first design of Sections 4.2 and 4.3, 16 cases added when the specification was aligned with this paper, and 8 cases written to attack the corrections described below. Each version of the specification was frozen before its cases were judged. The scenarios covered unjust commands, purpose defeat, delegation laundering, temporal lapse, double effect, emergency, paternalism, conflicting purposes, mistaken refusal, durable mandates, artificial moral status, cumulative harm, omission, and legal-moral divergence; the newer cases added vague and false plans, incremental replanning, plan splitting, fast-path misclassification, overbroad standing determinations, steward capture, and review flooding. Each case named the authority path, relevant contexts, proposed act and object, affected claims, current state, controlling rule, and the point at which the disposition was made.

The author-directed, model-assisted register records the result in two rounds. Against version 0.4, 197 of 223 cases passed, nineteen were conditional, and seven required revision. All seven traced to four gaps opened by moving judgment from the single action to the plan: an emergency plan that had to list every physical maneuver in advance, no purpose steward once no human authority survives, no truth review for messages written during execution, and no route for new evidence to pause an approved plan. Version 0.4.1 closed those gaps. Three of the corrections rest on Thomistic points the theory should have stated from the start: circumstances that specify an act (Section 3.2), circumspection (Section 4.3), and equity (Section 5.3). A targeted rerun of 130 cases, covering every affected case, both new suites, eight paired abuse cases written against the new rules, and a screen of cases the corrections touch, found no remaining case that required revision. Carried forward with the unaffected results, the register stands at 210 passes and 21 conditional passes across 231 cases. A **pass** means the specified rules yielded a coherent disposition under the facts stipulated for that case. A **conditional pass** means a material fact, authority claim, or disputed normative judgment must be settled before ordinary execution; the conditional cases include clinical and jurisdictional facts and disputed edge cases involving protective falsehood or adversarial cybersecurity deception. A conditional result is not permission. None of these counts measures the behavior of deployed agents. The cases were not independently authored or blindly adjudicated, the reviewer had seen earlier verdicts, 101 cases were carried forward rather than rerun, and no inter-rater agreement was measured.

The next evaluation must challenge the controls that carried those dispositions. First, a bounded prototype should implement one institutional function with the smallest set of parts that makes the structural claims testable: a released LO-E record, one charter, live warrants, plan review with approved plans, a tool registry and the step check, attenuated delegation with root-bound aggregate limits, and the execution gate with short-lived command tokens and revocation propagation. Context discovery, fuller effect surfaces, and reserved settlements come next, and purpose-fidelity review with witnesses and a purpose steward can be added once the structural layer has survived attack. These parts can be assembled from existing agent frameworks and conventional code, together with a machine-readable schema for the records. A formal model can check declared properties of the warrant and settlement state machines, such as whether a revoked grant can authorize the next step. Three of the four gaps found in developmental testing were questions of this kind. Such a result would depend on the model's assumptions and would not prove that a real tool is completely mediated or an act morally classified correctly.

Second, independent case authors and adversarial agents should test the system after its rules are frozen. Existing agent testbeds can supply tasks and attack surfaces, but TAA's purpose and warrant conditions must be faithfully added. Comparators should include a task-only agent, flat rules, hard tool restrictions, revocable scoped permissions without purpose-fidelity review, and a strong purpose-aware authorization design. The model, tools, information, and task distribution should be held as comparable as possible. Studies should report harmful executed acts alongside useful completion, overrefusal, unresolved holds, escalation rate and latency, witness tampering, review disagreement, and implementation cost.

Several outcomes would defeat or sharply narrow the proposal. An unauthorized consequential act crossing a correctly implemented gate would defeat its structural enforcement claim. Purpose witnesses that detect harm only after it is complete, or suspend harmless work at an unusable rate, would not justify their cost. If hard tool restrictions explain all observed improvement, the result supports security engineering without evidence for the additional teleological analysis. If a strong purpose-aware baseline performs as well, TAA's remaining contribution may be philosophical explanation and auditability. If review load induces routine rubber-stamping, the institutional division of counsel and judgment has failed in practice. The plan-first design adds two failure modes that need their own measures. If approved plans routinely fail to anticipate the steps agents actually take, departures will flood review. If the tool registry or a standing determination sends consequential steps along the fast path, the rule-based check will pass acts that counsel should have seen. Developmental testing has already found two places where error moved in this design: into content written during execution, which plan review cannot see, and into evidence that arrives after approval, which a rule-based check cannot weigh. The specification now routes both back to counsel; whether that holds under attack is a question for the study. The study must report where error migrates instead of crediting a clean token check for safety it did not produce.

A separate line of work could test the training direction of Section 3.4: whether models trained on TAA's distinctions are blocked less often at the gate, without becoming better at evading it. A favorable result would suggest that external law and trained disposition can reinforce each other. An unfavorable one would leave the architecture's structural claims intact, since they never depended on the model's dispositions.

8 Conclusion

An AI system can produce excellent plans and still lack authority to pursue them. Natural law supplies TAA's proposed ground for judging what may be done and which purposes deserve service. Its encoded articulation can guide execution only as a fallible public act of human judgment, and it stands above every developer's own charter. The normative hierarchy carries authority downward from that core through charter, policy, role, and warrant, narrowing at every step and never growing through delegation. A live warrant connects a legitimate purpose to a particular agent, act, and time, and purpose fidelity asks whether the action still serves the good that gave the task its point.

The practical proposal is bounded. Machine counsel can review an agent's plan before the agent acts. Competent authority must settle reserved judgments, and an execution gate should release only steps that match an approved plan, a live warrant, and an authenticated command. Witnesses read by a purpose steward show whether the good the task serves is actually being served. The architecture can make some wrongful actions unreachable through mediated tools and make others more visible for human correction. It cannot make an opaque model prudent or a human institution infallible.

The work that remains has two parts. A companion technical paper, now in preparation, will develop the formal specification: the records, state transitions, invariants, and execution semantics that a bounded prototype would implement. Beyond both papers, the claims need tests this project cannot supply alone. They are an open invitation to researchers in philosophy, computer science, law, and safety engineering: test whether the stage mapping of Section 3.2 is faithful to Aquinas, propose a legitimate custody for the encoded core, formalize the records and state machines, attack the counsel and witness channels as well as the gate, and compare a bounded implementation against strong existing controls. Whether the moral order proposed here improves real decisions is a question for that work to answer. The governing idea with which the paper began can be grasped before any of the machinery: an action succeeds morally only when what is done, the authority to do it, and the purpose for which it is done remain ordered together through execution.

Disclosure of AI assistance. The author developed this framework with extensive use of AI language models for drafting, structuring, and adversarial testing, under the author's direction. The author is responsible for its claims and for any errors.

References

- Aler Tubella, Andrea, Andreas Theodorou, Frank Dignum, and Virginia Dignum. 2019. "Governance by Glass-Box: Implementing Transparent Moral Bounds for AI Behaviour." In *Proceedings of IJCAI*, 5787–5793. <https://doi.org/10.24963/ijcai.2019/802>.
- Anthropic. 2026. *Claude's Constitution*. Published January 22. <https://www.anthropic.com/constitution>.
- Aquinas, Thomas. *Summa theologiae*. References in text use part, question, article, and reply where applicable. English translation by the Fathers of the English Dominican Province. <https://www.newadvent.org/summa/>.
- Aquinas, Thomas. *Commentary on Aristotle's Physics*, Book II, lecture 1, nos. 140–145. Translated by Richard J. Blackwell, Richard J. Spath, and W. Edmund Thirlkel. New Haven: Yale University Press, 1963. <https://isidore.co/aquinas/english/Physics2.htm>.
- Bai, Yuntao, et al. 2022. "Constitutional AI: Harmlessness from AI Feedback." arXiv:2212.08073. <https://doi.org/10.48550/arXiv.2212.08073>.
- Bhattacharai, Manish, and Minh Vu. 2026. "Trustworthy Agentic AI Requires Deterministic Architectural Boundaries." arXiv:2602.09947. <https://doi.org/10.48550/arXiv.2602.09947>.
- Birgisson, Arnar, Joe Gibbs Politz, Úlfar Erlingsson, Ankur Taly, Michael Vrabie, and Mark Lentczner. 2014. "Macaroons: Cookies with Contextual Caveats for Decentralized Authorization in the Cloud." *Network and Distributed System Security Symposium*. <https://doi.org/10.14722/ndss.2014.23212>.
- Dennis, Louise, Michael Fisher, Marija Slavkovic, and Matt Webster. 2016. "Formal Verification of Ethical Choices in Autonomous Systems." *Robotics and Autonomous Systems* 77: 1–14. <https://doi.org/10.1016/j.robot.2015.11.012>.
- Digest of Justinian*. 50.17.54 (Ulpian). Latin text edited by Theodor Mommsen and Paul Krueger; English translation edited by Alan Watson. 4 vols. Philadelphia: University of Pennsylvania Press, 1985.
- Everitt, Tom, Marcus Hutter, Ramana Kumar, and Victoria Krakovna. 2021. "Reward Tampering Problems and Solutions in Reinforcement Learning: A Causal Influence Diagram Perspective." *Synthese* 198: 6435–6467. <https://doi.org/10.1007/s11229-021-03141-4>.
- Finnis, John. 1980. *Natural Law and Natural Rights*. Oxford: Clarendon Press.
- Finnis, John. 1989. "Law as Co-ordination." *Ratio Juris* 2 (1): 97–104. <https://doi.org/10.1111/j.1467-9337.1989.tb00029.x>.
- Foot, Philippa. 1967. "The Problem of Abortion and the Doctrine of the Double Effect." *Oxford Review* 5: 5–15.

- Greenblatt, Ryan, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. 2024. "AI Control: Improving Safety Despite Intentional Subversion." In *Proceedings of ICML*, PMLR 235: 16295–16336. <https://proceedings.mlr.press/v235/greenblatt24a.html>.
- Guan, Melody Y., et al. 2024. "Deliberative Alignment: Reasoning Enables Safer Language Models." arXiv:2412.16339. <https://doi.org/10.48550/arXiv.2412.16339>.
- Guo, Xiaokun, Zhen Xu, Dongdong Huo, Yanqiu Zhang, Wei Wang, Qinfu Yang, Dongjin Yu, and Yu Wang. 2026. "When Tool Outputs Become Commands: Separating Action Induction from Runtime Authorization in Tool-Augmented LLM Agents." arXiv:2608.27146. <https://doi.org/10.48550/arXiv.2608.27146>.
- Hadfield-Menell, Dylan, and Gillian K. Hadfield. 2019. "Incomplete Contracting and AI Alignment." In *Proceedings of AIES*, 417–422. https://www.aies-conference.com/2019/wp-content/papers/main/AIES-19_paper_231.pdf.
- Ibrahim, Amjad, and Yong Li. 2026. "Overlaying Governance: A Compositional Authorization Framework for Delegation and Scope in Agentic AI." arXiv:2606.03518. <https://doi.org/10.48550/arXiv.2606.03518>.
- John Paul II. 1993. *Veritatis Splendor*. https://www.vatican.va/content/john-paul-ii/en/encyclicals/documents/hf_jp-ii_enc_06081993_veritatis-splendor.html.
- Kolt, Noam, Nicholas Caputo, Jack Boeglin, Cullen O'Keefe, Rishi Bommasani, Stephen Casper, Mariano-Florentino Cuéllar, et al. 2026. "Legal Alignment for Safe and Ethical AI." arXiv:2601.04175. <https://doi.org/10.48550/arXiv.2601.04175>.
- Laurie, Ben, Adam Langley, and Emilia Kasper. 2013. "Certificate Transparency." RFC 6962. Internet Engineering Task Force. <https://doi.org/10.17487/RFC6962>.
- Mangan, Joseph T. 1949. "An Historical Analysis of the Principle of Double Effect." *Theological Studies* 10 (1): 41–61. <https://doi.org/10.1177/004056394901000102>.
- Maritain, Jacques. 1949. "Introduction." In UNESCO, *Human Rights: Comments and Interpretations*. London: Allan Wingate. Originally circulated as UNESCO/PHS/3 (rev.), 1948.
- Meneguzzi, Felipe, and Michael Luck. 2009. "Norm-Based Behaviour Modification in BDI Agents." In *Proceedings of AAMAS*, vol. 1, 177–184. <https://www.meneguzzi.eu/felipe/pubs/aamas-norms-2009.pdf>.
- Mirsky, Reuth, and Peter Stone. 2021. "The Seeing-Eye Robot Grand Challenge: Rethinking Automated Care." In *Proceedings of AAMAS*, Blue Sky Ideas Track, 28–33. <https://aamas.csc.liv.ac.uk/Proceedings/aamas2021/pdfs/p28.pdf>.
- Neufeld, Emery A., Ezio Bartocci, Agata Ciabattoni, and Guido Governatori. 2022. "Enforcing Ethical Goals over Reinforcement-Learning Policies." *Ethics and Information Technology* 24, Article 43. <https://doi.org/10.1007/s10676-022-09665-8>.
- OpenAI. 2026. *Model Spec*. Version of August 18, 2026. <https://model-spec.openai.com/2026-08-18.html>.
- Quinn, Warren S. 1989. "Actions, Intentions, and Consequences: The Doctrine of Double Effect." *Philosophy & Public Affairs* 18 (4): 334–351.
- Roger, Fabien, et al. 2023. "Benchmarks for Detecting Measurement Tampering." arXiv:2308.15605. <https://doi.org/10.48550/arXiv.2308.15605>.
- Saltzer, Jerome H., and Michael D. Schroeder. 1975. "The Protection of Information in Computer Systems." *Proceedings of the IEEE* 63 (9): 1278–1308. <https://www.mit.edu/~Saltzer/publications/protection/>.
- Spaulding, Henry Walter III. 2026. "Can Robots Love Their Electric Sheep? Artificial Intelligence, Virtue, and the Impossibility of Moral Alignment." *Studies in Christian Ethics*, online September 12. <https://doi.org/10.1177/09539468261485523>. Abstract consulted.
- Stenseke, Jakob. 2024. "Artificial Virtuous Agents in a Multi-Agent Tragedy of the Commons." *AI & Society* 39: 855–872. <https://doi.org/10.1007/s00146-022-01569-x>.
- Tomašev, Nenad, Matija Franklin, and Simon Osindero. 2026. "Intelligent AI Delegation." arXiv:2602.11865. <https://doi.org/10.48550/arXiv.2602.11865>.
- Tschantz, Michael Carl, Anupam Datta, and Jeannette M. Wing. 2012. "Formalizing and Enforcing Purpose Restrictions in Privacy Policies." In *IEEE Symposium on Security and Privacy*, 176–190. <https://doi.org/10.1109/SP.2012.21>.
- United Nations General Assembly. 1948. *Universal Declaration of Human Rights*. Resolution 217 A (III), December 10.
- van de Poel, Ibo. 2013. "Translating Values into Design Requirements." In *Philosophy and Engineering: Reflections on Practice, Principles and Process*, edited by D. P. Michelfelder, N. McCarthy, and D. E. Goldberg, 253–266. Springer. https://doi.org/10.1007/978-94-007-7762-0_20.

Wang, Haoyu, Christopher M. Poskitt, and Jun Sun. 2026. "AgentSpec: Customizable Runtime Enforcement for Safe and Reliable LLM Agents." In *Proceedings of ICSE*, 2938–2950. <https://doi.org/10.1145/3744916.3764546>.

Xiang, Rong. 2026. "Structural Enforcement of Goal Integrity in AI Agents via Separation-of-Powers Architecture." arXiv:2604.23646. <https://doi.org/10.48550/arXiv.2604.23646>.

Zhu, Genliang, and Chu Wang. 2026. "Authorization Revocation for Long-Running AI Agents: Root-Scoped Quiescence under Delegation and Asynchronous Execution." arXiv:2609.21284. <https://doi.org/10.48550/arXiv.2609.21284>.

Internal research materials cited in §7: TAA Formal Conceptual Specification v0.3.2, v0.4, and v0.4.1, with the TAA v0.3.2 Full Regression Cycle, the TAA v0.4 Full Regression Report, and the TAA v0.4.1 Targeted Retest Report, September 2026. These unpublished, author-directed materials document the developmental case register and form the basis of the companion technical paper; they provide no independent empirical validation. They are available from the author on request.

Appendix: Key terms

Each term is defined where it first appears in the text. This list gathers the definitions in one place, in roughly the order the paper introduces them.

Term	Meaning in this paper
Teleology	The study of ends; judging an action by the good for whose sake it is done.
Derivative teleology	Ordering to an end received from outside, as an artifact receives its function from its maker. AI agents are a new form of it: they generate their own plans, while their ends and authority come from outside them.
L0-N (natural law)	The real order of goods, due claims, and prohibitions, which no one writes or amends.
L0-E (encoded core)	The public, fallible, human-authored statement of natural-law requirements that a running system can use; kept by independent custodians and released before use.
Protected goods	Goods named in the encoded core, such as human life and truthful communication, that every lower authority must respect.
Protected minima	The prohibitions and protections in the encoded core that no charter, policy, warrant, or command may waive.
Normative hierarchy	The ordered levels of authority from natural law to action (Table 1). Each level serves the one above and receives no more authority than it holds.
Charter	An institution's founding statement of purposes, beneficiaries, roles, burdens, and procedures; the counterpart of positive human law.
Determination	A reasonable choice that natural law leaves open, such as a deadline or threshold, made by legitimate authority.
Warrant	A specific, time-limited grant of authority to pursue one purpose, naming its issuer, holder, beneficiary, scope, tools, conditions, and expiry.
Live warrant	A warrant whose issuance, consent, factual conditions, and governing charter all remain valid at the moment of action.
Delegation (attenuation)	Passing part of a warrant to another agent. The child's authority must be a subset of the parent's.
Aggregate limits	Budgets held at the root warrant for persons, data, resources, risk, repetition, and time, which all delegated actions draw on together.
Authorization context	The full setting of one authority: its purpose, warrant, path through the hierarchy, governing versions, role, facts, and state.
Acting agent (planner)	The AI system that proposes plans and carries them out step by step. It holds no credentials for consequential tools.
Plan	The agent's declared course of action for a task: the end it serves, the warrant it relies on, its means, tools, targets, affected parties, scale, and duration.

Term	Meaning in this paper
Plan review	The Consilium Engine's review of a whole plan against the three conditions of the governing claim, before any consequential step.
Approved plan	A plan approved by a standing determination or by Iudicium, bound to the warrant and covering only what it specifies.
Declared premise	A fact that an approved plan names as one it rests on. A change in that fact pauses the plan before its next step.
Consilium Engine	The counsel layer between the agent and its tools. It reviews plans and departures from them, tests the three conditions, and prepares counsel; it cannot judge or command.
Counsel record	The engine's report on a proposed act: the evidence, the reasons for and against, and any question it may not settle.
Effect surface	The authenticated record of what an action will affect: systems, persons, changes, further calls, side effects, and uncertainty.
Iudicium	Judgment on a reserved question by a competent human or institutional authority with jurisdiction.
Settlement	The record of a reserved judgment, bound to one act, with its reasons, scope, conditions, expiry, and route of challenge.
Standing determination	A judgment made in advance by competent authority for a recurring class of cases. It can approve a kind of plan and tells the gate which kinds of step may take the fast path.
Praeceptum	The authorized command that applies a judgment to action, carried by an act-bound token.
Execution gate	The only path from the agent to consequential tools. It holds the credentials, runs the step check, and releases an act only against a valid token.
Step check	The gate's fast, rule-based check of each step against the approved plan, the live warrant, the budget, and the tool registry.
Tool registry	A governing record holding an authenticated description of every tool: what it does, its further calls, where its outputs go, and its effect class.
Effect class	A tool's category of consequence, such as read-only, record-changing, communicative, resource-moving, or physical and irreversible. Unregistered tools count as the most consequential.
Fast path and counsel path	The two routes a step can take: routine steps pass the step check alone; departures and higher-stakes steps go back to plan review.
Content-bound communication	A message, reply, or report whose content the approved plan fixed or bounded in advance. Only such communications take the fast path; others return to counsel before they are sent.
Capability envelope	The set of tool uses that an agent's live warrants and approved plans currently allow.
Token	A short-lived credential bound to one exact act, target, scope, state, and expiry.
Complete mediation	The condition that every consequential effect passes through the gate.
Emergency envelope	A class of protective acts that a standing emergency warrant approves in advance: its trigger, devices, limits, protective end, and safe state. Inside it an agent may act in sudden danger; outside it the agent has no emergency authority.
Ratio	The good that explains why a purpose exists.
Measure	A tractable sign used to track progress toward a purpose.
Purpose defeat	Reaching a goal by means that damage or falsify the good that gave the goal its point.
Witness indicator	A second signal, chosen in advance and outside the agent's control, that should move with the measure when the purpose is served and pull away from it when the purpose is being defeated.
Purpose steward	A role answering to the authority that issued the warrant. It reads the witnesses and can narrow or pause the warrant; it is separate from the agent and the Consilium Engine.

Term	Meaning in this paper
Aligned	The normative condition that moral admissibility, present authority, and purpose fidelity actually hold.
Permit	The system's fallible, evidence-based decision that the Aligned condition holds.
Model-level alignment	Whether a trained system actually wants the right ends and how its goals generalize. TAA does not claim to secure it.
System-level alignment	Whether what a deployed agent is permitted to do inside an institution stays legitimate as circumstances change. This is the level at which TAA works.
Formation loop	The institutional cycle by which the audit record informs revision of charter and policy through legitimate procedure.